

How it works in java. ConcurrentHashMap.



Sergey Kuptsov · [Follow](#)

11 min read · Apr 5, 2017



202



2



The main principle of programming says not to reinvent the wheel. But sometimes in order to understand what is going on and how not to misuse instrument we need to do this. Today reinventing concurrent hash map.

At first we need 2 things. I start with 2 tests — one indicates that our concurrent map implementation has no data races (actually we need to check if our test correct also by testing incorrect implementation) and second one that we will use to test performance in terms of throughput under different profiles of load.

Let's implement only some number of methods from Map interface:

```
public interface Map<K, V> {  
    V put(K key, V value);  
    V get(Object key);  
    V remove(Object key);  
    int size();  
}
```

Thread-safety correctness test

It is practically impossible to write thread safety test exhaustively, you need to take into consideration all aspects defined in [Chapter 17 of the JLS](#), more over it heavily depends on hardware [memory models](#) or JVM implementation.

For concurrency test let's use one of stress testing frameworks such as [jstress](#) that will heavy run your code trying to find data inconsistency. Although [jstress](#) is still experimental it is better choice. Why it is hard to write your own concurrency test — look at Shipilev's [lecture](#).

Let's start exploring jstress with [jstress-gradle-plugin](#). Full source code can be found in [how-it-works-concurrent-map](#).

```

public class ConcurrentMapThreadSafetyTest {

    @State
    public static class MapState {
        final Map<String, Integer> map = new HashMap<>(3);
    }

    @JCStressTest
    @Description("Test race map get and put")
    @Outcome(id = "0, 1", expect = ACCEPTABLE, desc = "return 0L and 1L")
    @Outcome(expect = FORBIDDEN, desc = "Case violating atomicity.")
    public static class MapPutGetTest {

        @Actor
        public void actor1(MapState state, LongResult2 result) {
            state.map.put("A", 0);
            Integer r = state.map.get("A");
            result.r1 = (r == null ? -1 : r);
        }

        @Actor
        public void actor2(MapState state, LongResult2 result) {
            state.map.put("B", 1);
            Integer r = state.map.get("B");
            result.r2 = (r == null ? -1 : r);
        }
    }

    @JCStressTest
    @Description("Test race map check size")
    @Outcome(id = "2", expect = ACCEPTABLE, desc = "size of map = 2")
    @Outcome(id = "1", expect = FORBIDDEN, desc = "size of map = 1 is race")
    @Outcome(expect = FORBIDDEN, desc = "Case violating atomicity.")
    public static class MapSizeTest {

        @Actor
        public void actor1(MapState state) {
            state.map.put("A", 0);
        }

        @Actor
        public void actor2(MapState state) {
            state.map.put("B", 0);
        }

        @Arbiter
        public void arbiter(MapState state, IntResult1 result) {
            result.r1 = state.map.size();
        }
    }
}

```

In first test MapPutGetTest we have two threads executing concurrently methods actor1 and actor2 respectively, both of them put some value to map, and checking them back, if there is no data race, both threads must see setted values.

In second MapSizeTest we concurrently put some different keys to map and after all checking the size of map — if there is no data race — the result must be 2.

We must check if this test correct both on non-thread-safe HashMap — we must observe atomicity violation and thread-safe ConcurrentHashMap — we must not see alternative results.

1. Results with non-thread-safe HashMap

```
[FAILED]
ru.skuptsov.concurrent.map.test.ConcurrentMapTest.MapPutGetTest
  Observed state  Occurrences  Expectation  Interpretation
        -1, 1      293,867    FORBIDDEN   Case violating atomic
        0, -1      282,190    FORBIDDEN   Case violating atomic
        0, 1      28,013,763  ACCEPTABLE  return 0 and 1

[FAILED]
ru.skuptsov.concurrent.map.test.ConcurrentMapTest.MapSizeTest
  Observed state  Occurrences  Expectation  Interpretation
        1         1,434,783  FORBIDDEN   size of map = 1 race
        2         11,733,097  ACCEPTABLE  size of map = 2
```

In not thread-safe HashMap we see some statistical number of inconsistent result, all 2 test failed

2. Results with thread-safe ConcurrentHashMap

```
[OK] ru.skuptsov.concurrent.map.test.ConcurrentMapTest.MapPutGetTest
  Observed state  Occurrences  Expectation  Interpretation
        0, 1      20,195,000    ACCEPTABLE

[OK] ru.skuptsov.concurrent.map.test.ConcurrentMapTest.MapSizeTest
  Observed state  Occurrences  Expectation  Interpretation
        2         6,573,730    ACCEPTABLE  size of map = 2
```

ConcurrentHashMap passed test, at least we can admit than our test can detect some simple concurrency issues. Same results for Collection.synchronizedMap and Hashtable.

First concurrent hash map attempt

The first naive approach is to simply synchronize every access to internal structures of map(array of buckets).

Actually we can write some concurrent wrapper over given Map provider — so does java.util.Collections#synchronizedMap, Hashtable and guava’s SynchronizedMultimap.

```
public class GeneralMonitorSynchronizedHashMap<K, V> extends
BaseMap<K, V> implements Map<K, V>, IMap<K, V> {

    private final Map<K, V> provider;
    private final Object monitor;

    public SynchronizedHashMap(Map<K, V> provider) {
        this.provider = provider;
        monitor = this;
    }

    @Override
    public V put(K key, V value) {
        synchronized (monitor) {
            return provider.put(key, value);
        }
    }
}
```

```

    }

    @Override
    public V get(Object key) {
        synchronized (monitor) {
            return provider.get(key);
        }
    }

    @Override
    public int size() {
        synchronized (monitor) {
            return provider.size();
        }
    }
}

```

Changes to non-volatile provider will be visible between threads, according to documentation:

Second, when a synchronized method exits, it automatically establishes a happens-before relationship with any subsequent invocation of a synchronized method for the same object. This guarantees that changes to the state of the object are visible to all threads.

Our SynchronizedHashMap passes concurrent tests but at what cost? in every method there can be only one thread at the same time even if we work with different keys, so in multithreaded load we must expect some performance penalty. Let's measure it.

Performance benchmark test

For performance test we will use [jmh](#). In performance benchmark w'll do not test non-thread-safe implementations.

```

@State(Scope.Thread)
@Warmup(iterations = 5, time = 1, timeUnit = TimeUnit.SECONDS)
@Measurement(iterations = 5, time = 1, timeUnit = TimeUnit.SECONDS)
@Fork(3)
@BenchmarkMode(Mode.AverageTime)
@OutputTimeUnit(MICROSECONDS)
public class ConcurrentMapBenchmark {
    private Map<Integer, Integer> map;

    @Param({"concurrenthashmap", "hashtable", "synchronizedhashmap"})
    private String type;

    @Param({"1", "10"})
    private Integer writersNum;

    @Param({"1", "10"})
    private Integer readersNum;

    private final static int NUM = 1000;

    @Setup
    public void setup() {
        switch (type) {
            case "hashtable":
                map = new Hashtable<>();
                break;
            case "concurrenthashmap":
                map = new ConcurrentHashMap<>();
                break;
        }
    }
}

```

```

        case "synchronizedhashmap":
            map = new SynchronizedHashMap<>(new HashMap<>());
            break;
    }
}

@Benchmark
public void test(Blackhole bh) throws ExecutionException,
InterruptedException {

    List<CompletableFuture> futures = new ArrayList<>();

    for (int i = 0; i < writersNum; i++) {
        futures.add(CompletableFuture.runAsync(() -> {
            for (int j = 0; j < NUM; j++) {
                map.put(j, j);
            }
        }));
    }

    for (int i = 0; i < readersNum; i++) {
        futures.add(CompletableFuture.runAsync(() -> {
            for (int j = 0; j < NUM; j++) {
                bh.consume(map.get(j));
            }
        }));
    }

    CompletableFuture.allOf(futures.toArray(new
    CompletableFuture[1])).get();
}
}

```

Benchmark	(readersNum)	(type)	(writersNum)	Mode	Cnt	Score	Error	Units
ConcurrentMapBenchmark.test	1	concurrenthashmap	1	avgt	15	65,157 ±	3,636	us/op
ConcurrentMapBenchmark.test	1	concurrenthashmap	10	avgt	15	302,023 ±	11,893	us/op
ConcurrentMapBenchmark.test	1	hashtable	1	avgt	15	156,149 ±	8,133	us/op
ConcurrentMapBenchmark.test	1	hashtable	10	avgt	15	736,649 ±	31,852	us/op
ConcurrentMapBenchmark.test	1	synchronizedhashmap	1	avgt	15	178,888 ±	11,110	us/op
ConcurrentMapBenchmark.test	1	synchronizedhashmap	10	avgt	15	1019,344 ±	67,259	us/op
ConcurrentMapBenchmark.test	10	concurrenthashmap	1	avgt	15	235,409 ±	6,956	us/op
ConcurrentMapBenchmark.test	10	concurrenthashmap	10	avgt	15	395,036 ±	23,586	us/op
ConcurrentMapBenchmark.test	10	hashtable	1	avgt	15	873,091 ±	65,774	us/op
ConcurrentMapBenchmark.test	10	hashtable	10	avgt	15	1564,516 ±	126,142	us/op
ConcurrentMapBenchmark.test	10	synchronizedhashmap	1	avgt	15	983,234 ±	109,709	us/op
ConcurrentMapBenchmark.test	10	synchronizedhashmap	10	avgt	15	1703,271 ±	162,345	us/op

We checked that performance of our SynchronizedHashMap is pretty much the same at java's Hashtable and it is 2 times worse than ConcurrentHashMap. Let's try to increase it.

Lock striping concurrent hash map attempt

The first improvement can come from the idea that instead of blocking access to the whole map it's better block thread access only to the same bucket where bucket array index = key.hashCode() % array.length. This technique is called lock striping or fine-grained synchronization, more other techniques here [«The Art of Multiprocessor Programming»](#).

For array of buckets we need an array of locks, at start the size of locks array must be equal to internal array size — it is important because we don't want situation where 2 locks are responsible for one array slot.

For simplicity we will design not resizable Map — that means that we cannot extend initial capacity(if $N \gg \text{initialCapacity}$ we will loose $O(1)$ map

gurance. Also we do not need loadFactor) — resizable concurrent map is separate big topic.

```
public class LockStripingArrayConcurrentHashMap<K, V> extends
BaseMap<K, V> implements Map<K, V> {

    private final AtomicInteger count = new AtomicInteger(0);
    private final Node<K, V>[] buckets;
    private final Object[] locks;

    @SuppressWarnings({"rawtypes", "unchecked"})
    public LockStripingArrayConcurrentHashMap(int capacity) {
        locks = new Object[capacity];
        for (int i = 0; i < locks.length; i++) {
            locks[i] = new Object();
        }

        buckets = (Node<K, V>[]) new Node[capacity];
    }

    @Override
    public int size() {
        return count.get();
    }

    @Override
    public V get(Object key) {
        if (key == null) throw new IllegalArgumentException();
        int hash = hash(key);
        synchronized (getLockFor(hash)) {
            Node<K, V> node = buckets[getBucketIndex(hash)];

            while (node != null) {
                if (isKeyEquals(key, hash, node)) {
                    return node.value;
                }

                node = node.next;
            }

            return null;
        }
    }

    @Override
    public V put(K key, V value) {
        if (key == null || value == null) throw new
IllegalArgumentException();
        int hash = hash(key);
        synchronized (getLockFor(hash)) {
            int bucketIndex = getBucketIndex(hash);
            Node<K, V> node = buckets[bucketIndex];

            if (node == null) {
                buckets[bucketIndex] = new Node<>(hash, key, value,
null);

                count.incrementAndGet();
                return null;
            } else {
                Node<K, V> prevNode = node;
                while (node != null) {
                    if (isKeyEquals(key, hash, node)) {
                        V prevValue = node.value;
                        node.value = value;

                        return prevValue;
                    }

                    prevNode = node;
                    node = node.next;
                }

                prevNode.next = new Node<>(hash, key, value, null);
                count.incrementAndGet();
                return null;
            }
        }
    }
}
```

```

    }
}

private boolean isKeyEquals(Object key, int hash, Node<K, V>
node) {
    return node.hash == hash &&
           node.key == key ||
           (node.key != null && node.key.equals(key));
}

private int hash(Object key) {
    return key.hashCode();
}

private int getBucketIndex(int hash) {
    return hash % buckets.length;
}

private Object getLockFor(int hash) {
    return locks[hash % locks.length];
}

private static class Node<K, V> {
    final int hash;
    K key;
    V value;
    Node<K, V> next;

    Node(int hash, K key, V value, Node<K, V> next) {
        this.hash = hash;
        this.key = key;
        this.value = value;
        this.next = next;
    }
}
}

```

It is important that all out class fields are final cause only final and static fields guarantee safe publication through constructor.

Source code can be found [here](#). Benchmark results:

Benchmark	{readersNum}	(type)	{writersNum}	Mode	Cnt	Score	Error	Units
ConcurrentMapBenchmark.test	1	generalmonitorsynchronizedmap	1	avgt	15	160,729 ±	12,324	us/op
ConcurrentMapBenchmark.test	1	generalmonitorsynchronizedmap	10	avgt	15	839,790 ±	39,961	us/op
ConcurrentMapBenchmark.test	1	lockarrayconcurrentmap	1	avgt	15	64,870 ±	2,617	us/op
ConcurrentMapBenchmark.test	1	lockarrayconcurrentmap	10	avgt	15	260,704 ±	7,499	us/op
ConcurrentMapBenchmark.test	1	concurrenthashmap	1	avgt	15	61,418 ±	2,167	us/op
ConcurrentMapBenchmark.test	1	concurrenthashmap	10	avgt	15	257,519 ±	5,234	us/op
ConcurrentMapBenchmark.test	10	generalmonitorsynchronizedmap	1	avgt	15	786,729 ±	89,060	us/op
ConcurrentMapBenchmark.test	10	generalmonitorsynchronizedmap	10	avgt	15	1469,810 ±	113,969	us/op
ConcurrentMapBenchmark.test	10	lockarrayconcurrentmap	1	avgt	15	320,662 ±	33,359	us/op
ConcurrentMapBenchmark.test	10	lockarrayconcurrentmap	10	avgt	15	482,189 ±	11,381	us/op
ConcurrentMapBenchmark.test	10	concurrenthashmap	1	avgt	15	219,265 ±	1,323	us/op
ConcurrentMapBenchmark.test	10	concurrenthashmap	10	avgt	15	381,790 ±	21,076	us/op

We can see that fine-grained synchronized implementation is better than our overall lock. Results for when there are one reader and one writer comparing with concurrent hash map are practically the same but when number of threads increases — the difference is bigger, especially where there are a lot of readers.

Lock free concurrent hash map attempt

Frankly speaking, synchronizing is not a parrallel programming technique cause it sets up threads in serial queue to wait for another thread to complete. And additional system cost of synchronization context switching

increasing as high as number of waiting threads grows but all we want is to make small number of instructions to change map's key value.

Let's rule some requirements to new hash map implementation that will improve our realization. And requirements are:

1. If we 2 threads that work with different keys(write or read) we do not want any kind of synchronization between them(cause word tearing is impossible in java — and access to two diff array fields is safe)
2. If multiple threads work on the same key(write and read) we do not want cache interleave(more about cache structure) and need safe happens-before guarantees for access between threads otherwise one thread might not see changed value by other thread. But we do not want to block read thread and wait for write thread to complete .
3. We do not want to block multiple readers to the same key if there is no one writer among them.

Let's concentrate on item 2 and 3. Actually we can make map read operation fully lock-free if we can make (1) volatile read array of buckets and then traverse inside bucket by linked list with (2) volatile read of next Node and value of Node itself which.

For (2) we can just mark Node's next and value fields volatile.

For the (1) there is no such thing as volatile array, even if it is declared as volatile, it not provides volatile semantics when reading or writing elements, concurrent accessing the k-th element of the array requires an explicit volatile read, volatile is only link to array. We can use AtomicReferenceArray for this purpose but it accepts only Object[] arrays. As alternative we can use Unsafe for volatile array read and lock-free write. The same technique is used in AtomicReferenceArray and ConcurrentHashMap.

```
@SuppressWarnings("unchecked")
// read array value by index
private <K, V> Node<K, V> volatileGetNode(int i) {
    return (Node<K, V>) U.getObjectVolatile(buckets, ((long) i <<
    ASHIFT) + ABASE);
}

// cas set array value by index
private <K, V> boolean compareAndSwapNode(int i, Node<K, V>
expectedNode, Node<K, V> setNode) {
    return U.compareAndSwapObject(buckets, ((long) i << ASHIFT) +
    ABASE, expectedNode, setNode);
}

private static final sun.misc.Unsafe U;
// Node[] header shift
private static final long ABASE;
// Node.class size shift
private static final int ASHIFT;

static {
    try {
```



```

// get unsafe by reflection - it is illegal to use not in java lib
Constructor<Unsafe> unsafeConstructor =
Unsafe.class.getDeclaredConstructor();
    unsafeConstructor.setAccessible(true);
    U = unsafeConstructor.newInstance();
} catch (NoSuchMethodException | InstantiationException |
InvocationTargetException | IllegalAccessException e) {
    throw new RuntimeException(e);
}

Class<?> ak = Node[].class;

ABASE = U.arrayBaseOffset(ak);
int scale = U.arrayIndexScale(ak);
ASHIFT = 31 - Integer.numberOfLeadingZeros(scale);
}

```

In volatileGetNode we can now read values safely with memory barrier.

Let's now write lock-free map V get(Object key) method:

```

public V get(Object key) {
    if (key == null) throw new IllegalArgumentException();
    int hash = hash(key);
    Node<K, V> node;

    // volatile read of bucket head at hash index
    if ((node = volatileGetNode(getBucketIndex(hash))) != null) {
        // check first node
        if (isKeyEquals(key, hash, node)) {
            return node.value;
        }

        // walk through the rest to find target node
        while ((node = node.next) != null) {
            if (isKeyEquals(key, hash, node))
                return node.value;
        }
    }

    return null;
}

```

We can check if we are really lock-free, according to [this](#).

In our first attempt we had big memory overhead with locks pool array — actually we can use the same fine-grained with additional memory — just lock on first node if it is exists. If it does not exists — we cannot block and need some lock-free method to set head of buckets — we already mentioned it above — method compareAndSwapNode

```

@Override
public V put(K key, V value) {
    if (key == null || value == null) throw new
IllegalArgumentException();
    int hash = hash(key);
    // no resize in this implementation - so the index will not
change
    int bucketIndex = getBucketIndex(hash);

    // cas loop trying not to miss
    while (true) {
        Node<K, V> node;

```

```

        // if bucket is empty try to set new head with cas
        if ((node = volatileGetNode(bucketIndex)) == null) {
            if (compareAndSwapNode(bucketIndex, null,
                new Node<>(hash, key, value, null))) {
                // if we succeed to set head - then break and return
                null

                count.increment();
                break;
            }
        } else {
            // head is not null - try to find place to insert or
            update under lock
            synchronized (node) {
                // check if node have not been changed since we got
                it

                // otherwise let's go to another loop iteration
                if (volatileGetNode(bucketIndex) == node) {
                    V prevValue = null;
                    Node<K, V> n = node;
                    while (true) {
                        ... simply walk through list under lock and
                        update or insert value...
                    }

                    return prevValue;
                }
            }
        }
    }
}

return null;
}

```

Full source code [here](#).

Let's benchmark it

Benchmark	(readersNum)	(type)	(writersNum)	Mode	Cnt	Score	Error	Units
ConcurrentMapBenchmark.test	1	lockarrayconcurrentmap	1	avgt	15	79,213 ±	2,966	us/op
ConcurrentMapBenchmark.test	1	lockarrayconcurrentmap	10	avgt	15	224,040 ±	11,307	us/op
ConcurrentMapBenchmark.test	1	lockfreearrayconcurrenthashmap	1	avgt	15	48,225 ±	0,619	us/op
ConcurrentMapBenchmark.test	1	lockfreearrayconcurrenthashmap	10	avgt	15	214,466 ±	4,785	us/op
ConcurrentMapBenchmark.test	1	concurrenthashmap	1	avgt	15	59,509 ±	1,260	us/op
ConcurrentMapBenchmark.test	1	concurrenthashmap	10	avgt	15	260,995 ±	6,053	us/op
ConcurrentMapBenchmark.test	10	lockarrayconcurrentmap	1	avgt	15	287,222 ±	6,790	us/op
ConcurrentMapBenchmark.test	10	lockarrayconcurrentmap	10	avgt	15	449,078 ±	3,857	us/op
ConcurrentMapBenchmark.test	10	lockfreearrayconcurrenthashmap	1	avgt	15	211,021 ±	7,620	us/op
ConcurrentMapBenchmark.test	10	lockfreearrayconcurrenthashmap	10	avgt	15	353,849 ±	19,990	us/op
ConcurrentMapBenchmark.test	10	concurrenthashmap	1	avgt	15	200,057 ±	3,749	us/op
ConcurrentMapBenchmark.test	10	concurrenthashmap	10	avgt	15	371,435 ±	10,971	us/op

We are even better than ConcurrentHashMap in some cases — but it is not a fair competition —cause ConcurrentHashMap do lazy table initialization during load and at least one resize cause resize number threshold = initialCapacity * loadFactor. If we run test again with initialCapacity != N inserted elements (=N/6)— results are slightly different.

Benchmark	(readersNum)	(type)	(writersNum)	Mode	Cnt	Score	Error	Units
ConcurrentMapBenchmark.test	1	lockarrayconcurrentmap	1	avgt	15	74,573 ±	5,742	us/op
ConcurrentMapBenchmark.test	1	lockarrayconcurrentmap	10	avgt	15	265,506 ±	26,030	us/op
ConcurrentMapBenchmark.test	1	lockfreearrayconcurrenthashmap	1	avgt	15	62,908 ±	2,358	us/op
ConcurrentMapBenchmark.test	1	lockfreearrayconcurrenthashmap	10	avgt	15	289,886 ±	14,090	us/op
ConcurrentMapBenchmark.test	1	concurrenthashmap	1	avgt	15	59,126 ±	0,959	us/op
ConcurrentMapBenchmark.test	1	concurrenthashmap	10	avgt	15	258,140 ±	3,263	us/op
ConcurrentMapBenchmark.test	10	lockarrayconcurrentmap	1	avgt	15	304,517 ±	2,466	us/op
ConcurrentMapBenchmark.test	10	lockarrayconcurrentmap	10	avgt	15	461,418 ±	16,804	us/op
ConcurrentMapBenchmark.test	10	lockfreearrayconcurrenthashmap	1	avgt	15	226,705 ±	4,925	us/op
ConcurrentMapBenchmark.test	10	lockfreearrayconcurrenthashmap	10	avgt	15	401,740 ±	11,943	us/op
ConcurrentMapBenchmark.test	10	concurrenthashmap	1	avgt	15	234,051 ±	9,835	us/op
ConcurrentMapBenchmark.test	10	concurrenthashmap	10	avgt	15	394,806 ±	16,960	us/op

This is because in ConcurrentHashMap we made resize during test and get element by key spends less time walking through bucket linked list.

Frankly speaking what we received is not a completely non-locking data structure — so does ConcurrentHashMap — but all we need is just to have a lock-free linked list — but with resizing and concurrent modification it is not so easy task — read [here](#).

Original java 8 ConcurrentHashMap has a number of small improvements we did not mentioned such as

- 1. Lazy table initialization that minimizes footprint until first use
- 2. Concurrent resizing array of buckets
- 3. Element count is maintained using a specialization of LongAdder which is one is well under high contention
- 4. Special types of bucket nodes(since 1.8) — TreeBins if the length of bucket list grows more than TREEIFY_THRESHOLD = 8 — bin become balanced tree with worst case key search ($O(\log(N_{bucket_size}))$)

Needless to say that implementation of ConcurrentHashMap in java 1.8 was significantly changed since 1.7. In 1.7 it was an idea of Segments where number of segment equals concurrencyLevel. Java 8 represents with a single array.

Java

Concurrency

Threads



Written by Sergey Kuptsov

117 Followers · 20 Following

Follow

Responses (2)



What are your thoughts?

Respond



Ilya Kharlamov
over 6 years ago



`return locks[hash % locks.length];`

this is obviously incorrect

since hash can be negative this will return a negative index

should be

```
return locks[hash & (locks.length-1)];
```

 1

Reply

 Mohan Radhakrishnan
over 5 years ago

...

“Special types of bucket nodes(since 1.8) — TreeBins if the length of bucket list grows more than TREEIFY_THRESHOLD = 8 — bin become balanced tree with worst case key search (O(log(Nbucket_size)))”

Are to referring to a red-black tree or something.....


[Read More](#)



Reply

More from Sergey Kuptsov



 Sergey Kuptsov


How it works in java. CompletableFuture.

The main principle of programming says not to reinvent the wheel. But sometimes in orde...

Nov 18, 2017

...



 Sergey Kuptsov


How it works in java. Thread Pool.

The main principle of programming says not to reinvent the wheel. But sometimes in orde...

Mar 4, 2017

...

jamsBenchOne.testCloneStream	Mode	Cnt	Score	Err
jamsBenchOne.testCloneStream:gc.alloc.rate	avgt	15	688,855 ±	653,9
jamsBenchOne.testCloneStream:gc.alloc.rate.norm	avgt	15	429,139 ±	220,8
jamsBenchOne.testCloneStream:gc.churn.PS_Eden_Space	avgt	15	302845,632 ±	1015,3
jamsBenchOne.testCloneStream:gc.churn.PS_Eden_Space.norm	avgt	15	427,356 ±	219,8
jamsBenchOne.testCloneStream:gc.churn.PS_Survivor_Space	avgt	15	303710,953 ±	48542,0
jamsBenchOne.testCloneStream:gc.churn.PS_Survivor_Space.norm	avgt	15	80,133 ±	0,1
jamsBenchOne.testCloneStream:gc.count	avgt	15	88,124 ±	50,4
jamsBenchOne.testCloneStream:gc.time	avgt	15	124,000	
jamsBenchOne.testCloneStream:gc.time	avgt	15	137,000	
jamsBenchOne.testJavaStream	avgt	15	342,575 ±	160,4
jamsBenchOne.testJavaStream:gc.alloc.rate	avgt	15	286,396 ±	126,9
jamsBenchOne.testJavaStream:gc.alloc.rate.norm	avgt	15	129288,148 ±	552,2
jamsBenchOne.testJavaStream:gc.churn.PS_Eden_Space	avgt	15	291,381 ±	126,3
jamsBenchOne.testJavaStream:gc.churn.PS_Eden_Space.norm	avgt	15	132750,937 ±	13554,4
jamsBenchOne.testJavaStream:gc.churn.PS_Survivor_Space	avgt	15	0,044 ±	0,0
jamsBenchOne.testJavaStream:gc.churn.PS_Survivor_Space.norm	avgt	15	22,721 ±	16,9
jamsBenchOne.testJavaStream:gc.count	avgt	15	91,000	
jamsBenchOne.testJavaStream:gc.time	avgt	15	116,000	

 Sergey Kuptsov


How it works in java. Streams.

The main principle of programming says not to reinvent the wheel. But sometimes in orde...

Feb 25, 2018

...



 Sergey Kuptsov

GraphQL vs REST

In web service with 2 separate teams—front and back — we have a communication...


Jan 9, 2019

...

See all from Sergey Kuptsov

Recommended from Medium



 In Stackademic by Abhinav


Is Java Really Dead? The Surprising Truth Behind Its Undeate...

The world of programming is ever-evolving, with new languages, frameworks, and...

★ Dec 20, 2024

...

(Available in Java 17 & Higher versions)	
Features	Benefits
Sealed Classes/Interfaces	Ensure a controlled hierarchy by restricting which classes or interfaces can extend or implement a parent.
Record Classes	Eliminates boilerplate and provides immutability by default. Replace verbose subclasses with concise record declarations & concise syntax for getters, constructors, and more.
Switch Enhancements	Supports multi-label case statements (e.g., case A, B:). Uses when keyword and -> syntax for direct value assignment without break statements, supports various types of patterns including Record types, guarded pattern, and allows concise, exhaustive handling of options for better code readability and control.
Pattern Matching for instanceof	Eliminates the need for instanceof by automatically casting into a single step. It simplifies type-checking and reduces boilerplate code.
Text Blocks	Multi-line string literals, enclosed in """, for cleaner handling of large text.
Factory Methods for Collections	Provides List.of(), Set.of(), and Map.of() for immutable collection creation.

 javatechonline blog

Java Interview Coding Problems Using Java 17 vs. Java 8 Features

In this post, we will explore some important Java Coding Interview questions with their...

★ Dec 8, 2024


...

Lists



General Coding Knowledge

20 stories · 1849 saves



Staff picks

793 stories · 1553 saves



data science and AI

40 stories · 310 saves



 In AlgoMaster.io by Ashish Pratap Singh


How I Tricked My Brain to Be Addicted to Coding

The Dopamine Hack

★ Nov 15, 2024



Mono	Flux
0 or 1	0 to many
Single API/database call	Streaming data
<code>Mono<T></code>	<code>Flux<T></code>

 Shishir Kumar

Reactive Programming in Java: Mono and Flux with Spring Boot

Introduction

★ Sep 12, 2024




 Full Stack Developer

Spring doesn't recommend @Autowired anymore???

There are 3 ways we can inject dependencies and Field Injection is not recommended...

★ Jul 19, 2024



 Engineering Digest

Multithreading in Java

CPU

Aug 2, 2024



See more recommendations